



- [Home](#)
 - [Recipes](#)
 - [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)
-



- 1.1. Installation
 - 1.1.1. Binaries
 - 1.1.2. Homebrew
- 1.2. Connecting to a Cluster
- 1.3. The config dotfiles
- 2. Introduction
 - 2.1. Navigating the Shell
 - 2.2. Getting Help
 - 2.3. The Prompt explained
 - 2.4. Loading Data into the Shell
 - 2.4.1. Manual import
 - 2.5. Exporting Data from the Shell
- 3. `cb-env` and the Environment
 - 3.1. `cb-env managed`
 - 3.2. `cb-env cluster`
 - 3.3. `cb-env bucket/scope/collection`
 - 3.4. `cb-env project/capella-organization`
 - 3.5. `cb-env register`
 - 3.6. `cb-env llm`
 - 3.7. Per command execution environments
 - 3.7.1. The `--clusters` flag
 - 3.7.2. The `--bucket`, `--scope`, `--collection` flags
- 4. Couchbase Commands
 - 4.1. buckets
 - 4.1.1. `buckets`
 - 4.1.2. `buckets config`
 - 4.1.3. `buckets create`
 - 4.1.4. `buckets drop`
 - 4.1.5. `buckets flush`
 - 4.1.6. `buckets get`
 - 4.1.7. `buckets load-sample`
 - 4.1.8. `buckets update`
 - 4.2. scopes
 - 4.2.1. `scopes`
 - 4.2.2. `scopes create`
 - 4.2.3. `scopes drop`
 - 4.3. collections
 - 4.3.1. `collections`
 - 4.3.2. `collections create`
 - 4.3.3. `collections drop`
 - 4.4. doc



- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



- 4.5. [nodes](#)
 - 4.6. [subdoc get](#)
 - 4.7. [query](#)
 - 4.7.1. [query](#)
 - 4.7.2. [query advise](#)
 - 4.7.3. [query indexes](#)
 - 4.8. [vector](#)
 - 4.8.1. [vector search](#)
 - 4.8.2. [vector create-index](#)
 - 4.8.3. [vector enrich-doc](#)
 - 4.8.4. [vector enrich-text](#)
 - 4.9. [ask](#)
 - 4.10. [version](#)
5. Reference
- 5.1. [Config File Format](#)
 - 5.2. [Credentials File Format](#)
6. Release Notes
- 6.1. 1.0.0 - 2024-09-11
 - 6.2. 0.75.2 - 2023-04-10
 - 6.3. 0.75.1 - 2023-04-13
 - 6.4. 0.75.0 - 2023-02-09

Welcome to the Couchbase Shell `cbsh` documentation! If you are new, you will want to start with the quickstart section and then proceed with the introduction. If you are already familiar with the shell, feel free to jump right into the couchbase commands.

Note that while the project is maintained by Couchbase, it is not covered under the EE support contract. We are providing community support through this [bug tracker](#) (<https://github.com/couchbaselabs/couchbase-shell/issues>).

1. Quickstart

1.1. Installation

The current latest version is **1.0.0**.

There are a couple ways you can get access to `cbsh`.

1.1.1. Binaries

Download our pre-built binaries for your platform of choice:

- Linux: [cbsh-x86_64-unknown-linux-gnu.tar.gz](#)
(https://github.com/couchbaselabs/couchbase-shell/releases/download/v1.0.0/cbsh-x86_64-unknown-linux-gnu.tar.gz)

**Couchbase**

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

linux-gnu.tar.gz

win.zip)

- Windows: [cbsh-x86_64-pc-windows-msvc.zip](https://github.com/couchbaselabs/couchbase-shell/releases/download/v1.0.0/cbsh-x86_64-pc-windows-msvc.zip)
(https://github.com/couchbaselabs/couchbase-shell/releases/download/v1.0.0/cbsh-x86_64-pc-windows-msvc.zip)

Once you've downloaded the zip file, extract it and switch into the just created directory. The following example shows it for mac, but it works very similar if you are on linux (just align the commands with the file you just downloaded):

```
$ unzip cbsh-x86_64-apple-darwin.zip
$ ls
cbsh LICENSE LICENSE_AGREEMENT README.md
```

You can now run the cbsh binary:

```
> ./cbsh --version
The Couchbase Shell 1.0.0
```

TIP

If you are running a recent macOS release (i.e. 10.15.x), you'll likely see an error similar to "**cbsh**" was blocked from use because it is not from an identified developer. This is because our binaries are not yet signed. To run it nonetheless you need to either navigate to System Preferences → Security & Privacy and click Allow Anyway , or run sudo xattr -r -d com.apple.quarantine \$PWD/cbsh inside your terminal. Next time you run the binary you'll get another prompt but then it should run fine.

1.1.2. Homebrew

If running on macOS you can install via the [Homebrew](https://formulae.brew.sh/formula/couchbase-shell) (<https://formulae.brew.sh/formula/couchbase-shell>) formula:

```
$ brew install couchbase-shell
```

Then run as follows:

```
$ cbsh --version
1.0.0
```

1.2. Connecting to a Cluster

The first time that you run ./cbsh you will receive a prompt asking if you'd like to create a config file. If you choose yes then the shell will provide you with a series of prompts to provide information about your default cluster. If you choose no then it will try to connect to localhost using the Administrator username and the password password.

You can modify this through CLI arguments (see ./cbsh -h for more information).

Note: Unless you specify TLS settings then PLAIN authentication is used and your credentials are sent in plaintext.

**Couchbase**

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

←-tls

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

ault, credentials will sent



Administrator 🏠 default

>

Once in the shell, you can start to execute commands (see the introduction section for more information). As a quick sanity check, list the nodes in the cluster:

> nodes

#	cluster	hostname	status	services	version	
0	dev.local	127.0.0.1:8091	healthy	search,indexing,kv,query	8.0.0-1246-enterprise	x86_64

To start experimenting with data operations load some sample data onto the cluster:

Administrator 🏠 default

> buckets load-sample travel-sample

#	cluster	sample	status
0	local	travel-sample	success

Now you can try running N1QL queries using the query command.

Administrator 🏠 default

> query "SELECT * FROM `travel-sample`.inventory.airline LIMIT 1"

#	airline	cluster
0	{id: 10, type: "airline", name: "40-Mile Air", iata: "Q5", icao: "MLA", callsign: "MILE-AIR", country: "United States"}	local

Or you can get documents by switching to the travel sample bucket with cb-env and using [doc_get](#) (https://couchbase.sh/docs/#_doc_get):



- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



				rror	cluster
0	airline_10				local
		id	10		
		type	airline		
		name	40-Mile Air		
		iata	Q5		
		icao	MLA		
		callsign	MILE-AIR		
		country	United States		

1.3. The config dotfiles

Connecting to a single cluster through the command line is nice when you are starting out, but later on you will likely either connect to the same cluster all the time or even to a multitude of them. To help with this, you can create a config file to hold your cluster details that the shell will read on startup.

The config file must be called `config` and be placed in a `.cbsh` dot file either in your home directory or in the directory from which the shell is being run. If you want to change the path of the directory where the config file is held this can be done with the `config-dir` flag when the shell is run:

```
> ls ~/config_file
config
> ./cbsh --config-dir ~/config_file
```

Note that even when the path to the directory containing the config file is given using this flag, the file containing the cluster information must still be called `config`.

The downloaded zip contains an example already, but here is a small sample config to help you get started as well:



- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



```
password = "password"
```

```
[[cluster]]
identifier = "remote"
connstr = "10.143.200.101"
default-bucket = "myapp"
username = "user"
password = "pass"
capella-organization = "my-org"

[[capella-organization]]
identifier = "my-org"
access-key = "get-your-own"
secret-key = "get-your-own"
default-project = "default"
```

This will register two clusters, one called `local` and one called `remote`. Now when you start the shell, it will connect to `local` automatically and you are all set. Changing between registered clusters is done using `cb-env cluster`.

The config also registers a Capella organization which is associated with the cluster "remote" using the "capella-organization" field, see `cb-env organizations` for more details. A config file can be created with just a capella-organization then the shell can be used to create and register a cluster for use.

Please check out the reference section on additional parameters you can set as well as how to move the credentials to a separate `credentials` file in case you want to share your config with other people and they do not use the same credentials.

2. Introduction

Couchbase Shell is fully featured, so it does not only contain commands related to Couchbase but is actually built on top of a general purpose shell called [Nushell](https://www.nushell.sh/) (<https://www.nushell.sh/>). This allows you to interact with the file system or any other command available on your machine, making it a great tool for both operational and development tasks on top of Couchbase.

The following introduction only touches on the basic concepts to make you productive quickly. We recommend also checking out the great [Nushell documentation](https://www.nushell.sh/book) (<https://www.nushell.sh/book>) so you can get the most out of it.

2.1. Navigating the Shell

Commands take inputs and produce output in a structured manner, most often represented as tables. Note how both the generic `ls` command and the Couchbase-specific `buckets` command both produce a table as their output:



Couchbase

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



3 Dockerfile	file	590 B	2 months ago	
4 LICENSE	file	11.1 KiB	2 months ago	
5 LICENSE AGREEMENT	file	43.7 KiB	2 months ago	
6 NOTICES	file	221.6 KiB	a week ago	
7 README.md	file	9.4 KiB	2 months ago	
8 build	dir	64 B	4 weeks ago	
9 docs	dir	608 B	3 minutes ago	
10 examples	dir	192 B	4 months ago	
11 jupyter	dir	128 B	4 months ago	
12 rust-toolchain.toml	file	524 B	2 weeks ago	
13 src	dir	416 B	a day ago	
14 target	dir	352 B	2 weeks ago	
15 tests	dir	160 B	4 months ago	

> buckets

#	cluster	name	type	replicas	min_durability_level	ram_quota	flush_enabled
0	default	beer-sample	couchbase	1	none	200.0 MiB	false
1	default	default	couchbase	2	none	100.0 MiB	true
2	default	travel-sample	couchbase	2	none	200.0 MiB	false

You can pipe the output into other commands, for example if you only want to see buckets that have `sample` in their name you can utilize the `where` command:

> buckets | where name =~ "sample"

#	cluster	name	type	replicas	min_durability_level	ram_quota	flush_enabled
0	default	beer-sample	couchbase	1	none	200.0 MiB	false
1	default	travel-sample	couchbase	2	none	200.0 MiB	false

In a similar fashion you can turn this structured table into other output formats, for example JSON:



- [Home](#)
- [Recipes](#)
- [Github](#) (<https://github.com/couchbaselabs/couchbase-shell>)



```

    "min_durability_level": "none",
    "ram_quota": 209715200,
    "flush_enabled": false,
    "status": "",
    "cloud": false
},
{
  "cluster": "default",
  "name": "travel-sample",
  "type": "couchbase",
  "replicas": 2,
  "min_durability_level": "none",
  "ram_quota": 209715200,
  "flush_enabled": false,
  "status": "",
  "cloud": false
}
]

```

Exactly this type of composition takes the unix philosophy of "do one thing well" and meshes it together with the idea of flexible structured pipelines. This allows to build powerful compositions that help you in your daily operations with Couchbase, both from a developer or operations point of view.

2.2. Getting Help

Other than using this documentation for help, each command can be called with `-h` or `--help` to get information about potential flags, arguments and subcommands. Also, some commands provide additional examples.

```
> buckets -h
Perform bucket management operations
```

Usage:

```
> buckets {flags}
```

Subcommands:

```
buckets config - Shows the bucket config (low level)
buckets create - Creates a bucket
buckets drop - Drops buckets through the HTTP API
buckets flush - Flushes buckets through the HTTP API
buckets get - Fetches buckets through the HTTP API
buckets load-sample - Load a sample bucket
buckets update - Updates a bucket
```

Flags:

```
-h, --help - Display the help message for this command
--clusters <String> - the clusters which should be contacted
```

Some commands (like the one above) only act as groupings for subcommands, like `from`, `to` or `doc`. Since they do not serve a purpose on their own, they will render their help output automatically:



- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



```
doc insert - Insert a document through the data service
doc remove - Removes a document through the data service
doc replace - Replace a document through the data service
doc upsert - Upsert (insert or override) a document through the data service
```

Flags:

-h, --help - Display the help message for this command

2.3. The Prompt explained

Couchbase Shell uses a custom, two line prompt to show you exactly in what environment you are working in right now. Since you can connect to different clusters, switch buckets etc. it is important to know what is currently "active", as the active buckets/scope/collection will be what commands are run against. Here is a sample prompt that will greet you when starting the shell:

```
👤 Administrator 🏠 local in 📁 travel-sample._default._default
>
```

It tells you that your user is `Administrator`, the current active cluster identifier is `local` and the active bucket is `travel-sample`. The current scope and collection have not been set, so they assume the `_default` value.

If you have an active scope or collection set then the prompt will also update to reflect that:

```
👤 Administrator 🏠 local in 📁 travel-sample.inventory.landmark
>
```

If we ran a `doc get` it would fetch the doc from `travel-sample.inventory.landmark`. In the second line, your actual user prompt starts.

2.4. Loading Data into the Shell

If you want to import data into Couchbase, or just load it into the shell for further processing, there are different commands available to help you. Once data is loaded into the shell the simplest way to save it to the connected cluster is using the `doc import` command. Alternatively if you want more control over the format and key of the uploaded data you can upload it manually using one of the other Couchbase save commands: `doc upsert` and `doc insert`.

2.4.1. Manual import

The `open` command will look at file endings and try to decode it automatically. Imagine a file named `user.json` in your current directory.



Couchbase

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



```
> open user.json
```

name	Michael
age	32
height	180

As you can see, the `open` command automatically decoded the JSON document into the tabular format. If the filename would only be `user`, the import would look like this instead:

```
> open user
{
  "name": "Michael",
  "age": 32,
  "height": 180
}
```

If you are dealing with data that cannot be decoded automatically, you can use the various `from` subcommands to help with decoding. In our case we use `from json`:

```
> open user | from json
```

name	Michael
age	32
height	180

The `doc upsert` and `doc insert` commands require there to be only two fields/columns. By default, these two columns are named `id` and `content`, but these can be overridden with `--id-column` and `--content-column`. The `id` column will be used as the key for the document, while the `content` column will be the JSON stored using that key. To get `user.json` in a format that can be inserted into the connected Cluster we [wrap](https://www.nushell.sh/commands/docs/wrap.html) (<https://www.nushell.sh/commands/docs/wrap.html>) the entire document into a `content` column and then [insert](https://www.nushell.sh/commands/docs/insert.html) (<https://www.nushell.sh/commands/docs/insert.html>) the `id` that we want to use:

```
> open user.json | wrap content | insert id $in.content.name
```

content		
name	Michael	
age	32	
height	180	
id	Michael	

There are many other approaches to achieving this same result. With our data in the correct format we can then upsert:

**Couchbase**

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



See the Importing data recipes for more examples.

2.5. Exporting Data from the Shell

The export counterparts to `open` and `from`, are `save` and `to`. You can use both commands to take tabular data from the shell and store it in files of the needed target format.

Like `open`, `save` will try to discern the format from the file ending. The following example will load a JSON file and then save it as CSV:

```
> cat user.json
{
  "name": "Michael",
  "age": 32,
  "height": 180
}

> open user.json | save user.csv

> cat user.csv
name,age,height
Michael,32,180
```

This example is dealing with only one row for simplicity, but you can save as many rows as you need in one file.

As a motivating example, the following snippet runs a N1QL query and stores the result as a csv file:

```
> query "select airportname,city,country from `travel-sample` where type = 'airport' limit 10" | save output.csv

> cat output.csv
airportname,city,country
Calais Dunkerque,Calais,France
Peronne St Quentin,Peronne,France
Les Loges,Nangis,France
Couterne,Bagnole-de-l'orne,France
Bray,Albert,France
Le Touquet Paris Plage,Le Tourquet,France
Denain,Valenciennes,France
Glisy,Amiens,France
La Garenne,Agen,France
Cazaux,Cazaux,France
```

See the Exporting data recipes for more information.



Couchbase

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

host) active. The same is true for commands against (this can be



The `cb-env` command will tell you which resources are currently active (you are also able to tell from the prompt):

```
> cb-env
```

username	charlie
display_name	Charlie
cluster	capella
bucket	default
scope	inventory
collection	hotel
cluster_type	provisioned
capella-organization	couchbase
capella-project	Charlie work
llm	Gemini

If you were to now run a command then we would be running it:

- As the user "charlie"
- Against the "capella" cluster
- Against the "default" bucket
- Against the "inventory" scope
- Against the "hotel" collection

Note that `display_name` is the name that appears in your shell prompt and is not used by commands.

You can also change the active resources with the `cb-env` command.



- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



cb-env bucket - Sets the active bucket based on its name
 cb-env capella-organization - Sets the active capella organization based on its identifier
 cb-env cluster - Sets the active cluster based on its identifier
 cb-env collection - Sets the active collection based on its name
 cb-env managed - Lists all clusters currently managed by couchbase shell
 cb-env project - Sets the active project based on its name
 cb-env register - Registers a cluster for use with the shell
 cb-env scope - Sets the active scope based on its name
 cb-env timeouts - Sets the active timeouts for operations
 cb-env unregister - Unregisters a cluster for use with the shell

Flags:

-h, --help - Display the help message for this command
 --capella - show default execution environment of capella
 --timeouts - show default execution environment for timeouts

3.1. cb-env managed

Lists all the clusters you have registered with the shell.

```
> cb-env managed
```

#	active	tls	identifier	username	capella_organization
0	true	false	dev.local	Administrator	
1	false	true	capella	charlie	

3.2. cb-env cluster

Changes the active cluster. The change of cluster will be reflected in the output of cb-env and the prompt:

```
👤 Charlie 🏠 capella
```

```
> cb-env cluster local
```

```
👤 Charlie 🏠 local
```

```
> cb-env
```

username	charlie
display_name	Charlie
cluster	local
bucket	
scope	
collection	
cluster_type	other

Notice that when you are connected to a locally running (non-Capella) cluster then the capella specific fields (project and capella organization) will not appear.



- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

not registered:

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



and.

3.3. cb-env bucket/scope/collection

The cb-env bucket , cb-env scope and cb-env collection commands are used to set the active bucket/scope/collection respectively. When changing the bucket/scope/collection the change will be reflected in the prompt and the output of the cb-env command:

```
👤 Charlie 🏠 local
> cb-env bucket travel-sample
👤 Charlie 🏠 local in 🏫 travel-sample._default._default
> cb-env scope inventory
👤 Charlie 🏠 local in 🏫 travel-sample.inventory._default
> cb-env collection landmark
👤 Charlie 🏠 local in 🏫 travel-sample.inventory.landmark
>
```

Note that when we set an active bucket the scope and collection are set to the `_default` scope and collection.

When a "higher level" entity is changed the "lower level" entities are returned to this default setting:

```
👤 Charlie 🏠 local in 🏫 travel-sample.inventory.landmarks
> cb-env scope some_scope
👤 Charlie 🏠 local in 🏫 travel-sample.some_scope._default
> cb-env bucket some_bucket
👤 Charlie 🏠 local in 🏫 some_bucket._default._default
>
```

Setting the scope resets the collection and setting the bucket will reset both the scope and collection. It is important to remember that these commands do not validate the existence of resources when the environment is changed. Setting the active bucket/scope/collection to one that doesn't exist will only have an impact when you try and perform operations against it.

```
👤 Charlie 🏠 local
> cb-env bucket not-a-real-bucket
👤 Charlie 🏠 local in 🏫 not-a-real-bucket._default._default
> doc get test_doc
+-----+
| # | id | content | cas | error
+-----+
| 0 |    |        | 0 | Failed to load cluster config: bucket 'not-a-real-bucket' could not be found
+-----+
```

Checking that the active environment is correctly set to resources that exist is a good first port of call if things aren't working as expected.

**Couchbase**

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

tern contain clusters. Although you
; required to perform management



organization command:

```
> cb-env capella-organization my-org
```

Or if you switch to an active cluster that has a default Capella organization specified on registration then that org will become active. The logic is similar with the active project which can either be changed manually:

```
👤 Administrator 🏠 local in 🏢 default._default._default
> cb-env project my-project
```

Alternatively if the active organization has a default project defined on registration then that default project will become active. Similarly to changing active clusters, the `cb-env capella-organization` command will return an error if the named org has not been registered with the shell:

```
> cb-env capella-organization not-an-org
Error:   ✘ Organization not registered
        help: Has the organization not-an-org been registered in the config file?
```

But the same is not true with projects, so if operations are unexpectedly failing against a Capella cluster, make sure your active project is correctly set.

3.5. cb-env register

```
> cb-env register --help
Registers a cluster for use with the shell
```

Usage:

```
> cb-env register {flags} <identifier> <connstr> <username> <password>
```

Flags:

```
-h, --help - Display the help message for this command
--display_name <String> - the display name to use for the user when this cluster is active
--default-bucket <String> - the default bucket to use with this cluster
--default-scope <String> - the default scope to use with this cluster
--default-collection <String> - the default collection to use with this cluster
--tls-enabled <String> - whether or not to enable tls, defaults to true
--tls-cert-path <String> - the path to the certificate to use with tls
--tls-accept-all-certs <String> - whether or not to accept all certs with tls, defaults to true
--save - whether or not to add the cluster to the .cbsh config file, defaults to false
--capella-organization <String> - capella organization that this cluster belongs to
```

Parameters:

```
identifier <string>: the identifier to use for this cluster
connstr <string>: the connection string to use for this cluster
username <string>: the username to use for this cluster
password <string>: the password to use for this cluster
```

**Couchbase**

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



entry to the config file and restart.
provided for use. cb-env register
al parameters:

istor password

Once registered the new cluster can be seen in the output of `cb-env managed`, and can be set as the active cluster using `cb-env cluster`:

```
👤 Charlie 🏠 remote in ✅ default._default._default
> cb-env managed
```

#	active	tls	identifier	username	capella_organization
0	false	false	local	Administrator	
1	false	true	new-cluster	Administrator	
2	true	true	remote	charlie	my-org

Note that by default although tls is enabled the shell will accept all certs. This can be changed with the `--tls-accept-all-certs` flag.

3.6. cb-env llm

```
> cb-env llm --help
Sets the active llm based on its identifier
```

Usage:

```
> cb-env llm <identifier>
```

Flags:

```
-h, --help - Display the help message for this command
```

Parameters:

```
identifier <string>: the identifier of the llm
```

This command sets the active llm (large language model), which will be used by the vector enrich-doc, vector enrich-text and ask commands.

To be set as active the llm must be specified in the config file. For example:

**Couchbase**

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



[[llm]]

```
identifier = "Bedrock-titan"
provider = "Bedrock"
embed_model = "amazon.titan-embed-text-v1"
chat_model = "amazon.titan-text-express-v1"
```

[[llm]]

```
identifier = "Gemini-pro"
provider = "Gemini"
embed_model = "text-embedding-004"
chat_model = "gemini-1.0-pro"
api_key = "get-your-own"
```

The currently supported providers are Gemini (Google), Bedrock (AWS) and OpenAI. Specifying values other than these for the provider will result in an error when starting the shell. Notice that the Bedrock entry does not have an API key, this is because it requires the user to configure an appropriate role using the [AWS CLI](https://docs.aws.amazon.com/cli/v1/userguide/cli-configure-role.html) (<https://docs.aws.amazon.com/cli/v1/userguide/cli-configure-role.html>).

The `embed-model` field is the model that will be used to generate embeddings by the `vector enrich-doc` and `vector enrich-text` commands. While the `chat-model` is the model that will be used to answer questions with the `ask` command. These models can be any that the provider's API supports, and should be provided in the format given in the provider's API docs.

The api-keys can also be given separately in the credentials file, for example:

[[llm]]

```
identifier = "Gemini-pro"
api_key = "get-your-own"
```

The identifier must be the same as the entry in the config file for this to work.

The active llm can be checked using the `cb-env` command:

```
Charlie ~ remote in _default._default._default
> cb-env
+-----+
| username      | Administrator |
| display_name | Charlie      |
| cluster       | remote       |
| bucket        | default      |
| scope         | _default     |
| collection    | _default     |
| cluster_type  | provisioned |
| capella-organization | my-org      |
| llm           | Bedrock-titan |
+-----+
```

When the active llm is changed using `cb-env llm` this will be reflected in the output of `cb-env`:



- [Home](#)
- [Recipes](#)
- [Github](#) (<https://github.com/couchbaselabs/couchbase-shell>)



bucket	default
scope	_default
collection	_default
cluster_type	provisioned
capella-organization	my-org
llm	Gemini-pro

3.7. Per command execution environments

On many commands you will notice a set of flags which allow you to override the active execution environment. Different commands support different flags, depending on the command you can expect to see any of:

- `--clusters`
- `--bucket`
- `--scope`
- `--collection`

3.7.1. The `--clusters` flag

The argument for this flag is an identifier combined with a regular expression. So imagine you have three clusters setup with the following names:

```
> cb-env managed | get identifier
[{"id": 0, "name": "prod-us-east"}, {"id": 1, "name": "prod-us-west"}, {"id": 2, "name": "prod-eu-center"}]
```

If you wanted to run a command against all clusters in `prod-us`, you could use `--clusters prod-us.*`, e.g.

```
> buckets --clusters prod-us.*
```

#	cluster	name	type	replicas	min_durability_level	ram_quota	flush_e
0	prod-us-east	default	couchbase	1	none	200.0 MiB	false
1	prod-us-west	default	couchbase	1	none	200.0 MiB	false
2	prod-us-west	travel-sample	couchbase	1	none	200.0 MiB	false

In the background this gets passed to a regex engine, so you can go a little crazy with it if needed.

3.7.2. The `--bucket` , `--scope` , `--collection` flags

These flags are a little different to the `--clusters` flag, they are not regular expressions and can only be used to define a single name each. Unlike `--clusters` the name provided to these flags does not have to be already known to Couchbase Shell, they can refer to any bucket, scope, and collection that exist within your active cluster or defined cluster(s). For example:

**Couchbase**

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

users

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

171



	addresses	#	type	address	city	country
		0	home	0622 Adams Mills	Manchester	United Kingdom
	driving_licence			5f5f145d-a4db-5630-b7d8-874df29a505d		
	passport			a1c4f1ac-a7d7-5b97-88ed-11cafc634896		
	preferred_email			rigobertobernier@gadugca.sd		
	preferred_phone			(965) 227-3977		
	preferred_airline			inventory.airline.airline_5479		
	preferred_airport			inventory.airport.airport_478		
	credit_cards	#	type	number	expiration	
		0	American Express	346533746753899	2021-04	
	created			2020-04-12		
	updated			2021-02-19		

4. Couchbase Commands

The following sections discuss the individual Couchbase specific commands in greater detail. Remember, you can always mix and match them with built-in other shell commands as well as executable programs from your environment.

4.1. buckets

The `buckets` commands are used to perform bucket management operations. They can be particularly powerful combined with the `--clusters` (https://couchbase.sh/docs/#_the_clusters_flag) flag.

4.1.1. buckets

Lists all the buckets from your active cluster:

Charlie dev.local

> buckets

#	cluster	name	type	replicas	min_durability_level	ram_quota	flush_enabled
0	dev.local	beer-sample	couchbase	2	none	412.0 MiB	false
1	dev.local	default	couchbase	0	none	512.0 MiB	false
2	dev.local	memd	memcached	0	none	100.0 MiB	false

Check this [recipe](https://couchbase.sh/docs/recipes/#_moving_data_between_clusters) (https://couchbase.sh/docs/recipes/#_moving_data_between_clusters) to see how `cbsh` can be used to copy these buckets to another cluster.

4.1.2. buckets config

Retrieves the full config for the named bucket:



bucketType	membase
collectionsManifestUid	21
compressionMode	passive
conflictResolutionType	seqno
controllers	{record 4 fields}
ddocs	{record 1 field}
durabilityMinLevel	none
evictionPolicy	valueOnly
localRandomKeyUri	/pools/default/buckets/default/localRandomKey
maxTTL	0
name	default
nodeLocator	vbucket
nodes	[table 1 row]
numVBuckets	64
pitrEnabled	false
pitrGranularity	600
pitrMaxHistoryAge	86400
quota	{record 2 fields}
replicaIndex	false
replicaNumber	0
stats	{record 3 fields}
storageBackend	couchstore
streamingUri	/pools/default/bucketsStreaming/default?bucket_uuid=0ef162c33e14b163630f04639b347937
threadsNumber	3
uri	/pools/default/buckets/default?bucket_uuid=0ef162c33e14b163630f04639b347937
uuid	0ef162c33e14b163630f04639b347937
vBucketServerMap	{record 4 fields}

If you are unsure what you would use this for, you probably don't need it.

4.1.3. buckets create

Creates a bucket with the specified RAM quota on the active cluster:

👤 Charlie 🏠 local	> buckets create default 256
👤 Charlie 🏠 local	> buckets
# cluster name type replicas min_durability_level ram_quota flush_enabled cl	
0 local default couchbase 1 none 256.0 MiB false fc	

Check this [recipe](https://couchbase.sh/docs/recipes.html#_managing_multiple_clusters) (https://couchbase.sh/docs/recipes.html#_managing_multiple_clusters) to see how `cbsh` can help find a cluster for your bucket.

4.1.4. buckets drop

Drops the named bucket:

**Couchbase**

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



```

    ↗ Charlie 🏠 local
> buckets drop default
    ↗ Charlie 🏠 local
> buckets

```

empty list

ram_quota	flush_enabled	cl
56.0 MiB	false	fc

4.1.5. buckets flush

Deletes all docs from the named buckets:

```

    ↗ Charlie 🏠 local
> query "SELECT * FROM `travel-sample`" | length
31591
    ↗ Charlie 🏠 local
> buckets flush travel-sample
    ↗ Charlie 🏠 local
> query "SELECT * FROM `travel-sample`" | length
0

```

Not following what's going on, check the [query](#) (https://couchbase.sh/docs/#_query) and [length](#) (https://www.nushell.sh/commands/docs/length.html) documentation.

4.1.6. buckets get

Gets the named bucket:

```
> buckets get travel-sample
```

#	cluster	name	type	replicas	min_durability_level	ram_quota	flush_enabled
0	local	travel-sample	couchbase	1	none	200.0 MiB	true

4.1.7. buckets load-sample

Loads the named [sample bucket](#) (https://docs.couchbase.com/server/current/manage/manage-settings/install-sample-buckets.html) onto the active cluster:



- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



👤 Charlie 🏠 local

> buckets

#	cluster	name	type	replicas	min_durability_level	ram_quota	flush_enable
0	local	travel-sample	couchbase	1	none	200.0 MiB	false

4.1.8. buckets update

Updates the settings of an existing bucket:

👤 Charlie 🏠 local

> buckets

#	cluster	name	type	replicas	min_durability_level	ram_quota	flush_enable
0	local	travel-sample	couchbase	1	none	200.0 MiB	false

👤 Charlie 🏠 local

> buckets update travel-sample --flush true --replicas 2 --expiry 100

👤 Charlie 🏠 local

> buckets

#	cluster	name	type	replicas	min_durability_level	ram_quota	flush_enable
0	local	travel-sample	couchbase	2	none	200.0 MiB	true

4.2. scopes

The scopes commands are used to manage scopes.

4.2.1. scopes

Lists all of the scopes in the active bucket:

👤 Charlie 🏠 local in 📁 travel-sample._default._default

> scopes

#	scope	cluster
0	inventory	local
1	tenant_agent_00	local
2	tenant_agent_01	local
3	tenant_agent_02	local
4	tenant_agent_03	local
5	tenant_agent_04	local
6	_default	local



- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



#	scope	cluster
0	_default	local

```
👤 Charlie 🏠 local in ✐ default._default._default
> scopes create new-scope
👤 Charlie 🏠 local in ✐ default._default._default
> scopes
```

#	scope	cluster
0	new-scope	local
1	_default	local

4.2.3. scopes drop

Deletes the scope matching the given name:

```
👤 Charlie 🏠 local in ✐ default._default._default
> scopes
```

#	scope	cluster
0	new-scope	local
1	_default	local

```
👤 Charlie 🏠 local in ✐ default._default._default
> scopes drop new-scope
👤 Charlie 🏠 local in ✐ default._default._default
> scopes
```

#	scope	cluster
0	_default	local

4.3. collections

The `collections` commands are used to manage collections.

4.3.1. collections

Lists all of the collections in the active scope:



- [Home](#)
- [Recipes](#)
- [Github](#) (<https://github.com/couchbaselabs/couchbase-shell>)



2 airport	inherited	local	
3 airline	inherited	local	
4 route	inherited	local	

4.3.2. collections create

Create a collection with the name supplied:

```
👤 Charlie 🏠 local in 📁 my-bucket.my-scope._default
> collections
```

empty list

```
👤 Charlie 🏠 local in 📁 my-bucket.my-scope._default
> collections create new-collection
```

```
👤 Charlie 🏠 local in 📁 my-bucket.my-scope._default
> collections
```

#	collection	max_expiry	cluster
0	new-collection	inherited	local

4.3.3. collections drop

Drop the collection matching the name given:

```
👤 Charlie 🏠 local in 📁 default.my-scope._default
> collections
```

#	collection	max_expiry	cluster
0	new-collection	inherited	local

```
👤 Charlie 🏠 local in 📁 default.my-scope._default
> collections drop new-collection
```

```
👤 Charlie 🏠 local in 📁 default.my-scope._default
> collections
```

empty list

4.4. doc

The `doc` commands are for managing the documents stored in the registered clusters.

4.4.1. doc get

Gets a doc from the active cluster, bucket, scope and collection:

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



		type airline						
		name 40-Mile Air						
		iata Q5						
		icao MLA						
		callsign MILE-AIR						
		country United States						

The contents of the document are returned in the `content` column alongside the document metadata (id, cas, error and cluster).

`doc get` can get multiple documents at once using an input stream:

```
👤 Charlie 🏠 local in ⏱ travel-sample._default._default
> [airline_10 airline_10748 airline_137] | wrap id | doc get
```

#	id	content	cas	error	cluster
0	airline_10	id 10 type airline name 40-Mile Air iata Q5 icao MLA callsign MILE-AIR country United States	1712321628975398912		prod-us-west
1	airline_137	id 137 type airline name Air France iata AF icao AFR callsign AIRFRANS country France	1712321633323712512		prod-us-west
2	airline_10748	id 10748 type airline name Locair iata ZQ icao LOC callsign LOCAIR country United States	1712321631323947008		prod-us-west



Couchbase

- [Home](#)
- [Recipes](#)
- [Github](https://github.com/couchbaselabs/couchbase-shell) (<https://github.com/couchbaselabs/couchbase-shell>)

put column using the --id-

- [Home](#)
- [Recipes](#)
- [Github](https://github.com/couchbaselabs/couchbase-shell) (<https://github.com/couchbaselabs/couchbase-shell>)



supported by the [from](#)

(<https://www.nushell.sh/commands/docs/from.html>) command.

```
👤 Charlie 🏠 local in 🗂 default._default._default
> cat user.json
{
  "id": 123,
  "name": "Michael",
  "age": 32,
  "height": 180
}
```

```
👤 Charlie 🏠 local in 🗂 default._default._default
> doc import user.json
```

#	processed	success	failed	failures	cluster
0	1	1	0		local

```
👤 Charlie 🏠 local in 🗂 default._default._default
> doc get 123
```

#	id	content	cas	error	cluster								
0	123	<table border="1"> <tr> <td>id</td><td>123</td></tr> <tr> <td>name</td><td>Michael</td></tr> <tr> <td>age</td><td>32</td></tr> <tr> <td>height</td><td>180</td></tr> </table>	id	123	name	Michael	age	32	height	180	1726063078042501120		local
id	123												
name	Michael												
age	32												
height	180												

doc import will use the id field in the source document as the document key by default, but this behaviour can be changed with the --id-column flag.



- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



#	processed	success	failed	failures	cluster
0	1	1	0		local

👤 Charlie 🏠 local in ✂ default._default._default

> doc get Michael

#	id	content	cas	error	cluster						
0	Michael	<table border="1"> <tr> <td>name</td><td>Michael</td></tr> <tr> <td>age</td><td>32</td></tr> <tr> <td>height</td><td>180</td></tr> </table>	name	Michael	age	32	height	180	1726063937243250688		local
name	Michael										
age	32										
height	180										

TIP

look at the many different import formats `from` supports, including csv, xml, yaml and even sqlite. With this simple tool at hand you are able to load many different data formats quickly and import them into Couchbase!

4.4.3. doc insert

Inserts document into the active bucket/scope/collection:

👤 Charlie 🏠 remote in ✂ default._default._default

> open user.json

name	Michael
age	32
height	180

👤 Charlie 🏠 remote in ✂ default._default._default

> open user.json | wrap content | insert id \$in.content.name | doc insert

#	processed	success	failed	failures	cluster
0	1	1	0		remote

And if a document already exists in the active collection with this key then the command will fail:

👤 Charlie 🏠 remote in ✂ default._default._default

> open user.json | wrap content | insert id \$in.content.name | doc insert

#	processed	success	failed	failures	cluster
0	1	0	1	Key already exists	remote

- [Home](#)
- [Recipes](#)
- [Github \(<https://github.com/couchbaselabs/couchbase-shell>\)](#)

doc insert command.

»les.

- [Home](#)
- [Recipes](#)
- [Github \(<https://github.com/couchbaselabs/couchbase-shell>\)](#)



the following stored:

```
👤 Charlie 🏠 remote in ↴ default._default._default
```

```
> doc get Michael
```

#	id	content	cas	error	cluster						
0	Michael	<table border="1"> <tr> <td>name</td><td>Michael</td></tr> <tr> <td>age</td><td>32</td></tr> <tr> <td>height</td><td>180</td></tr> </table>	name	Michael	age	32	height	180	1726821299020365824		remote
name	Michael										
age	32										
height	180										

We would remove it as follows:

```
👤 Charlie 🏠 remote in ↴ default._default._default
```

```
> doc remove Michael
```

#	processed	success	failed	failures	cluster
0	1	1	0		remote

If a document matching the key cannot be found, then an error is returned:

```
👤 Charlie 🏠 remote in ↴ default._default._default
```

```
> doc remove Michael
```

#	processed	success	failed	failures	cluster
0	1	0	1	Key not found	remote

You can also remove multiple documents at once with an input stream:

```
👤 Charlie 🏠 remote in ↴ travel-sample._default._default
```

```
> [airline_10 airline_10748 airline_137] | wrap id | doc remove
```

#	processed	success	failed	failures	cluster
0	3	3	0		remote

4.4.5. doc replace

Replaces the document in Couchbase matching the key id of the new one, if there is no document matching the id then an error is returned. Say we have the following document stored in the connected cluster:



- [Home](#)
- [Recipes](#)
- [Github](#) (<https://github.com/couchbaselabs/couchbase-shell>)



			age	32			
			height	180			

And we have an updated version stored locally:

```
👤 Charlie 🏠 remote in `default._default._default
> open user.json
[{"name": "Michael", "age": 80, "height": 110}]
```

Then we can replace the first with the second using `doc replace`:

```
👤 Charlie 🏠 remote in `default._default._default
> open user.json | wrap content | insert id $in.content.name | doc replace
[{"#": 0, "processed": 1, "success": 1, "failed": 0, "failures": "", "cluster": "remote"}]
👤 Charlie 🏠 remote in `default._default._default
> doc get Michael
[{"#": 0, "id": "Michael", "content": {"name": "Michael", "age": 80, "height": 110}, "cas": "1726821910540648448", "error": "", "cluster": "remote"}]
```

See the manual import section for an explanation of the formatting we do before piping to the `doc replace` command. If there was no document with the `id` `Michael`, then the replace would fail:

```
👤 Charlie 🏠 remote in `default._default._default
> open user.json | wrap content | insert id $in.content.name | doc replace
[{"#": 0, "processed": 1, "success": 0, "failed": 1, "failures": "Key not found", "cluster": "remote"}]
```

Similarly to `doc insert`, `doc replace` can be used to replace multiple documents at once, see importing data for examples.



Couchbase

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](#)

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](#)



#	id	content	cas	error	cluster
0	Michael		0	Key not found	remote

👤 Charlie 🏠 remote in ↳ default._default._default
> open user.json

name	Michael
age	32
height	180

👤 Charlie 🏠 remote in ↳ default._default._default
> open user.json | wrap content | insert id \$in.content.name | doc upsert

#	processed	success	failed	failures	cluster
0	1	1	0		remote

👤 Charlie 🏠 remote in ↳ default._default._default
> doc get Michael

#	id	content	cas	error	cluster						
0	Michael	<table border="1"> <tr> <td>name</td> <td>Michael</td> </tr> <tr> <td>age</td> <td>32</td> </tr> <tr> <td>height</td> <td>180</td> </tr> </table>	name	Michael	age	32	height	180	1726822249315041280		remote
name	Michael										
age	32										
height	180										

See the manual import section for an explanation of the formatting we do before piping to the `doc upsert` command.

Or if there is an existing document with the same `id` then `upsert` will behave the same as a `replace`:



- [Home](#)
- [Recipes](#)
- [Github](#) (<https://github.com/couchbaselabs/couchbase-shell>)



			age	32			
			height	180			

👤 Charlie 🏠 remote in ↴ default._default._default

> open user.json

name	Michael
age	80
height	110

👤 Charlie 🏠 remote in ↴ default._default._default

> open user.json | wrap content | insert id \$in.content.name | doc upsert

#	processed	success	failed	failures	cluster
0	1	1	0	0	remote

👤 Charlie 🏠 remote in ↴ default._default._default

> doc get Michael

#	id	content	cas	error	cluster						
0	Michael	<table border="1"> <tr> <td>name</td><td>Michael</td></tr> <tr> <td>age</td><td>80</td></tr> <tr> <td>height</td><td>110</td></tr> </table>	name	Michael	age	80	height	110	1726822368839598080	0	remote
name	Michael										
age	80										
height	110										

4.5. nodes

The nodes command allows you to list all the nodes of the cluster you are currently connected to.

> nodes

#	cluster	hostname	status	services	version
0	prod-us-west	192.168.107.128:8091	healthy	search,indexing,kv,query	7.6.2-3505-enterprise
1	prod-us-west	192.168.107.129:8091	healthy	search,indexing,kv,query	7.6.2-3505-enterprise
2	prod-us-west	192.168.107.130:8091	healthy	search,indexing,kv,query	7.6.2-3505-enterprise

**Couchbase**

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

the data service

**Flags:**

-h, --help - Display the help message for this command
 --id-column <String> - the name of the id column if used with an input stream
 --bucket <String> - the name of the bucket
 --scope <String> - the name of the scope
 --collection <String> - the name of the collection
 --clusters <String> - the clusters which should be contacted
 --batch-size <Number> - the maximum number of items to batch send at a time
 -e, --halt-on-error - halt on any errors

Parameters:

path <any>: the path(s) to be fetched from the documents
 id <string>: the document id (optional)

It can be used to retrieve a field from a single document:

```
👤 Administrator 🏠 cluster in 📁 travel-sample._default._default
> subdoc get address landmark_10019
```

#	id	content	cas	error	cluster
0	landmark_10019	Prince Arthur Road, ME4 4UG	1722410659053961216		local

Or similarly to the doc commands a stream of ids can be provided:

```
👤 Administrator 🏠 cluster in 📁 travel-sample._default._default
> [landmark_10019 landmark_10020] | subdoc get address
```

#	id	content	cas	error	cluster
0	landmark_10019	Prince Arthur Road, ME4 4UG	1722410659053961216		local
1	landmark_10020	4 High Street, ME7 1BB	1722410654151999488		local

The path parameter can be a list, allowing retrieval of multiple fields in one or more docs:

**Couchbase**

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

cas | error |



0	landmark_10019	1722410659053961216	
local	name Royal Engineers Museum		
	address Prince Arthur Road, ME4 4UG		
1	landmark_10020	1722410654151999488	
local	name Hollywood Bowl		
	address 4 High Street, ME7 1BB		

4.7. query

The query commands can be used to explore/create indexes and execute queries.

4.7.1. query

Takes a N1QL statement and executes it against the active cluster.

```
👤 Charlie 🏠 local in 🏢 travel-sample._default._default
> query "SELECT meta().id FROM `travel-sample`.inventory.landmark WHERE country = 'France'"
```

#	id	cluster
0	landmark_10061	local
...
387	landmark_9838	local

The query gets all the IDs of the docs where the country is France, then `cbsh` appends the cluster column to the results.

Named parameters are supported through the `--params` flag when the argument is a json object:

- [Home](#)
- [Recipes](#)
- [Github](https://github.com/couchbaselabs/couchbase-shell) (<https://github.com/couchbaselabs/couchbase-shell>)



387	landmark_9838	local
-----	---------------	-------

Multiple named parameters can be used at once, note there is no need to separate them with commas:

```
👤 Charlie 🏠 local in 🏢 travel-sample._default._default
> query "SELECT airline FROM `travel-sample`.inventory.route WHERE sourceairport = $aval AND distance
  <table border='1'>
  <thead>
    <tr> <th>#</th> <th>airline</th> <th>cluster</th> </tr>
  </thead>
  <tbody>
    <tr> <td>0</td> <td>B6</td> <td>local</td> </tr>
    <tr> <td>1</td> <td>EK</td> <td>local</td> </tr>
    <tr> <td>2</td> <td>SV</td> <td>local</td> </tr>
  </tbody>
</table>
```

The wildcard character '%' can be used the same as inside of the query statement. The following finds any IDs that match the regex 'hotel1002*!.

```
👤 Charlie 🏠 local in 🏢 travel-sample._default._default
> query "SELECT meta().id FROM `travel-sample`.inventory.hotel WHERE meta().id LIKE $pattern" --params
  <table border='1'>
  <thead>
    <tr> <th>#</th> <th>id</th> <th>cluster</th> </tr>
  </thead>
  <tbody>
    <tr> <td>0</td> <td>hotel_10025</td> <td>local</td> </tr>
    <tr> <td>1</td> <td>hotel_10026</td> <td>local</td> </tr>
  </tbody>
</table>
```

Positional parameters are also supported when the `--params` flag takes a [list](https://www.nushell.sh/book/working_with_lists.html) (https://www.nushell.sh/book/working_with_lists.html) argument:

```
👤 Charlie 🏠 local
> query "SELECT airline FROM `travel-sample`.inventory.route WHERE sourceairport = $1 AND distance > $2"
  <table border='1'>
  <thead>
    <tr> <th>#</th> <th>airline</th> <th>cluster</th> </tr>
  </thead>
  <tbody>
    <tr> <td>0</td> <td>B6</td> <td>local</td> </tr>
    <tr> <td>1</td> <td>EK</td> <td>local</td> </tr>
    <tr> <td>2</td> <td>SV</td> <td>local</td> </tr>
  </tbody>
</table>
```

4.7.2. query advise

Helps you to learn about the indexes that your queries are using, and what indexes you could create to make them faster. For example we can take the first query from the `query` examples:

**Couchbase**

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

WHERE country = 'France'"

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



adviseinfo	current_indexes	#	0 CREATE PRIMARY INDEX def_inventory_landmark_primary ON `travel-sample`.`inventory`(`country`)
	recommended_indexes	covering_indexes	#
		indexes	#

This shows our query statement uses the Primary index `def_inventory_landmark_primary` and also recommends a covering index that we can create using `query`:

```
> query "CREATE INDEX adv_country ON `default`:`travel-sample`.`inventory`.`landmark`(`country`)"
```

4.7.3. query indexes

Lists all of the query indexes. If we had just created the recommended index at the end of the `query advise` section, the nushell command `find` can be used in conjunction with `query indexes` to check it was successfully created.

```
> query indexes | find adv_country
```

#	bucket	condition	index_key	keyspace	name	primary	scope	s
0	travel-sample		0 `country`	landmark	adv_country	false	inventory	c

And if we want to check the definition this can be done using the `--definitions` flag:

```
> query indexes --definitions | find adv_country
```

#	bucket	scope	collection	name	status	storage_mode	replicas	l
0	travel-sample	inventory	landmark	adv_country	Ready	plasma	0	CREATE

**Couchbase**

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

how query indexes can be

ita. Check the recipes for how the
[recipes.html#_similarity_search](#) or [simple](#)[RAG \(https://couchbase.sh/docs/recipes.html#_simple_rag\).](#)

4.8.1. vector search

Performs a vector search against a named vector index. Requires a vector to be used as the source of the search which can be supplied in a couple of formats.

For example imagine we have the following document stored in our cluster:

```
> doc get landmark_10019
```

#	id																																																						
0	landmark_10019																																																						
	<table border="1"> <tr> <td>title</td> <td>Gillingham (Kent)</td> </tr> <tr> <td>name</td> <td>Royal Engineers Museum</td> </tr> <tr> <td>alt</td> <td></td> </tr> <tr> <td>address</td> <td>Prince Arthur Road, ME4 4UG</td> </tr> <tr> <td>directions</td> <td></td> </tr> <tr> <td>phone</td> <td>+44 1634 822839</td> </tr> <tr> <td>tollfree</td> <td></td> </tr> <tr> <td>email</td> <td></td> </tr> <tr> <td>url</td> <td>http://www.remuseum.org.uk</td> </tr> <tr> <td>hours</td> <td>Tues - Fri 9.00am to 5.00pm, Sat - Sun 11.30am - 5.00pm</td> </tr> <tr> <td>image</td> <td></td> </tr> <tr> <td>price</td> <td></td> </tr> <tr> <td>content</td> <td>Adult - £6.99 for an Adult ticket that allows you to come back to the history of the British Empire. A quite extensive collection of tank mounted bridges etc can be seen for free. ticket.</td> </tr> <tr> <td>geo</td> <td> <table border="1"> <tr> <td>lat</td> <td>51.39</td> </tr> <tr> <td>lon</td> <td>0.54</td> </tr> <tr> <td>accuracy</td> <td>RANGE_INTERPOLATED</td> </tr> </table> </td> </tr> <tr> <td>activity</td> <td>see</td> </tr> <tr> <td>type</td> <td>landmark</td> </tr> <tr> <td>id</td> <td>10019</td> </tr> <tr> <td>country</td> <td>United Kingdom</td> </tr> <tr> <td>city</td> <td>Gillingham</td> </tr> <tr> <td>state</td> <td></td> </tr> <tr> <td>contentVector</td> <td> <table border="1"> <tr> <td>0</td> <td>0.02</td> </tr> <tr> <td>...</td> <td>...</td> </tr> <tr> <td>1023</td> <td>-0.00</td> </tr> </table> </td> </tr> </table>	title	Gillingham (Kent)	name	Royal Engineers Museum	alt		address	Prince Arthur Road, ME4 4UG	directions		phone	+44 1634 822839	tollfree		email		url	http://www.remuseum.org.uk	hours	Tues - Fri 9.00am to 5.00pm, Sat - Sun 11.30am - 5.00pm	image		price		content	Adult - £6.99 for an Adult ticket that allows you to come back to the history of the British Empire. A quite extensive collection of tank mounted bridges etc can be seen for free. ticket.	geo	<table border="1"> <tr> <td>lat</td> <td>51.39</td> </tr> <tr> <td>lon</td> <td>0.54</td> </tr> <tr> <td>accuracy</td> <td>RANGE_INTERPOLATED</td> </tr> </table>	lat	51.39	lon	0.54	accuracy	RANGE_INTERPOLATED	activity	see	type	landmark	id	10019	country	United Kingdom	city	Gillingham	state		contentVector	<table border="1"> <tr> <td>0</td> <td>0.02</td> </tr> <tr> <td>...</td> <td>...</td> </tr> <tr> <td>1023</td> <td>-0.00</td> </tr> </table>	0	0.02	1023	-0.00
title	Gillingham (Kent)																																																						
name	Royal Engineers Museum																																																						
alt																																																							
address	Prince Arthur Road, ME4 4UG																																																						
directions																																																							
phone	+44 1634 822839																																																						
tollfree																																																							
email																																																							
url	http://www.remuseum.org.uk																																																						
hours	Tues - Fri 9.00am to 5.00pm, Sat - Sun 11.30am - 5.00pm																																																						
image																																																							
price																																																							
content	Adult - £6.99 for an Adult ticket that allows you to come back to the history of the British Empire. A quite extensive collection of tank mounted bridges etc can be seen for free. ticket.																																																						
geo	<table border="1"> <tr> <td>lat</td> <td>51.39</td> </tr> <tr> <td>lon</td> <td>0.54</td> </tr> <tr> <td>accuracy</td> <td>RANGE_INTERPOLATED</td> </tr> </table>	lat	51.39	lon	0.54	accuracy	RANGE_INTERPOLATED																																																
lat	51.39																																																						
lon	0.54																																																						
accuracy	RANGE_INTERPOLATED																																																						
activity	see																																																						
type	landmark																																																						
id	10019																																																						
country	United Kingdom																																																						
city	Gillingham																																																						
state																																																							
contentVector	<table border="1"> <tr> <td>0</td> <td>0.02</td> </tr> <tr> <td>...</td> <td>...</td> </tr> <tr> <td>1023</td> <td>-0.00</td> </tr> </table>	0	0.02	1023	-0.00																																																
0	0.02																																																						
...	...																																																						
1023	-0.00																																																						



Couchbase

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

ut to a vector search. Due to the

[shell.sh/commands/docs/select.html](#))

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



#	contentVector
0	0 0.03

1023	-0.00

This can then be piped directly into vector search :

```
> doc get landmark_10019 | flatten | select contentVector | vector search landmark-content-index contentVector
```

#	id	score	cluster
0	landmark_10019	3402823500	local
1	landmark_28965	1.0286634	local
2	landmark_3547	1.0150017	local

We specify the vector index to be queried (landmark-content-index) as well as the name of the field on which the vector index was created (contentVector). See vector create-index for further explanation.

Subdoc get can also be used to fetch the source vector:

```
> subdoc get contentVector landmark_10019 | select content | vector search landmark-content-index contentVector
```

#	id	score	cluster
0	landmark_10019	3402823500	local
1	landmark_28965	1.0286634	local
2	landmark_3547	1.0150017	local

Here we don't need to flatten and specify the field since the content of the sub doc output already only holds the contentVector field.

Another format that is accepted is the output of vector enrich-text :

**Couchbase**

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

ark-content-index

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



	1		landmark_20732		0.7003826		local	
	2		landmark_21682		0.6987926		local	

Note the `--dimension` flag used here, the vector provided as the source of the search must match the dimension of the vectors on which the index was created else you will get no results.

Finally if you just have a plain vector that is a list of floats this can be passed in as a final positional parameter:

```
> let vector = [0.1 0.2 0.3 0.4]
> vector search vector-index fieldName $vector
```

The vector here only has a dimension of 4 which is unrealistically low, but it is used for illustrative purposes.

The output of the search is a table containing the IDs of the `n` (specified by the `--neighbours` flag) documents with the most similar vector in the field named. The similarity of vectors is calculated using the metric specified on index creation.

To see the contents of the documents identified in the table simply pipe the output of vector search into a [doc_get](#) (https://couchbase.sh/docs/#_doc_get) command. If only interested in a particular field then the subdoc get command will also work in the same way:

```
👤 Charlie 🏠 local in _default._default._default
> vector search landmark-content-index contentVector $vector | subdoc get address
+-----+
| # |      id      |          content          |      cas      | error | cluster |
+-----+
| 0 | landmark_11956 | Hornchurch Road, Hornchurch, RM11 1JU | 1722871671832641536 |      | local   |
| 1 | landmark_25284 | 4040 Twiggs St                         | 1722871685304025088 |      | local   |
| 2 | landmark_7744  | Grandstand Road                        | 1722871709487202304 |      | local   |
+-----+
```

The environment is important when using the `vector search` command as the fully qualified index name is constructed from the active bucket and scope. The above command will be trying to use the index `default._default.landmark-content-index`. If the index was created against a different bucket/scope then you will get an index not found error.

**Couchbase**

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



```
help: Unexpected status code: 400, body: {"error":"rest_auth: preparePerms, err: index not found","r
...
...
```

If you want to use an index that was created against a different bucket/scope to those that are active, this can be done with the `--bucket` and `--scope` flags.

```
👤 Charlie 🏠 local in 📁 travel-sample.inventory._default
> vector search landmark-content-index contentVector $vector --bucket default --scope _default
+----+-----+-----+-----+
| # |   id   |    score    | cluster |
+----+-----+-----+-----+
| 0 | landmark_10019 | 340282350000000000000000000000000000000000000000000000000000000000 | local   |
| 1 | landmark_16379 | 1.0082568   | local   |
| 2 | landmark_33857 | 0.9897698   | local   |
+----+-----+-----+-----+
```

4.8.2. vector create-index

Creates a vector index against the active bucket/scope or the bucket/scope specified with the corresponding flags. For example the following will create an index named new-vector-index over the field `fieldName` in documents in the bucket `default` and the scope `_default` :

```
👤 Charlie 🏠 local in ✎ default._default._default
> vector create-index new-vector-index fieldName 1024
```

Any documents located elsewhere will not be indexed, and a vector search performed with this index will not find them. If the documents to be indexed were stored elsewhere then the environment would need to be changed with `cb-env` bucket/scope or the bucket and scope could be passed in with the appropriate flags.

The `field` parameter is the name of the field in the documents that contains the vector, and the `dimension` (length of the vector) parameter must match that of the vectors contained in the named field. Imagine there are documents with the following structure in the bucket `users` and scope `embeddedUsers` :

```
{
  "name": "my-name",
  "e-mail": "shell-user@couchbase.com",
  "description": "A rich textual description of the user.",
  "descriptionEmbedding": [0.0178287, 0.123098, -0.0189239, 1.109238],
}
```

In this example the `descriptionEmbedding` is the result of using a large language model to generate an embedding from the `description` field. The dimension of the vector (4) is unrealistically low but it's been chosen for illustrative purposes. To create an index over such documents in order to find users with similar descriptions the command would be:



- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



cet users --scope

embeddedUsers

In the first example the create command is run from within the bucket and scope that the documents are stored in. Whereas in the latter the environment values need to be overwritten with the --bucket and --scope flags to correctly create the index.

4.8.3. vector enrich-doc

Enriches an existing json document by generating an embedding using the active ldm from a named field and storing it in a new field in the same document.

For example take one of the documents from the travel sample data set:

```
> doc get landmark_10019
```

#	id																																														
0	landmark_10019																																														
	<table border="1"> <tr> <td>title</td><td>Gillingham (Kent)</td></tr> <tr> <td>name</td><td>Royal Engineers Museum</td></tr> <tr> <td>alt</td><td></td></tr> <tr> <td>address</td><td>Prince Arthur Road, ME4 4UG</td></tr> <tr> <td>directions</td><td></td></tr> <tr> <td>phone</td><td>+44 1634 822839</td></tr> <tr> <td>tollfree</td><td></td></tr> <tr> <td>email</td><td></td></tr> <tr> <td>url</td><td>http://www.remuseum.org.uk</td></tr> <tr> <td>hours</td><td>Tues - Fri 9.00am to 5.00pm, Sat - Sun 11.30am - 5.00pm</td></tr> <tr> <td>image</td><td></td></tr> <tr> <td>price</td><td></td></tr> <tr> <td>content</td><td>Adult - £6.99 for an Adult ticket that allows you to come back f of the British Empire. A quite extensive collection that takes tank mounted bridges etc can be seen for free. There is also an</td></tr> <tr> <td>geo</td><td> <table border="1"> <tr> <td>lat</td><td>51.39</td></tr> <tr> <td>lon</td><td>0.54</td></tr> <tr> <td>accuracy</td><td>RANGE_INTERPOLATED</td></tr> </table> </td></tr> <tr> <td>activity</td><td>see</td></tr> <tr> <td>type</td><td>landmark</td></tr> <tr> <td>id</td><td>10019</td></tr> <tr> <td>country</td><td>United Kingdom</td></tr> <tr> <td>city</td><td>Gillingham</td></tr> <tr> <td>state</td><td></td></tr> </table>	title	Gillingham (Kent)	name	Royal Engineers Museum	alt		address	Prince Arthur Road, ME4 4UG	directions		phone	+44 1634 822839	tollfree		email		url	http://www.remuseum.org.uk	hours	Tues - Fri 9.00am to 5.00pm, Sat - Sun 11.30am - 5.00pm	image		price		content	Adult - £6.99 for an Adult ticket that allows you to come back f of the British Empire. A quite extensive collection that takes tank mounted bridges etc can be seen for free. There is also an	geo	<table border="1"> <tr> <td>lat</td><td>51.39</td></tr> <tr> <td>lon</td><td>0.54</td></tr> <tr> <td>accuracy</td><td>RANGE_INTERPOLATED</td></tr> </table>	lat	51.39	lon	0.54	accuracy	RANGE_INTERPOLATED	activity	see	type	landmark	id	10019	country	United Kingdom	city	Gillingham	state	
title	Gillingham (Kent)																																														
name	Royal Engineers Museum																																														
alt																																															
address	Prince Arthur Road, ME4 4UG																																														
directions																																															
phone	+44 1634 822839																																														
tollfree																																															
email																																															
url	http://www.remuseum.org.uk																																														
hours	Tues - Fri 9.00am to 5.00pm, Sat - Sun 11.30am - 5.00pm																																														
image																																															
price																																															
content	Adult - £6.99 for an Adult ticket that allows you to come back f of the British Empire. A quite extensive collection that takes tank mounted bridges etc can be seen for free. There is also an																																														
geo	<table border="1"> <tr> <td>lat</td><td>51.39</td></tr> <tr> <td>lon</td><td>0.54</td></tr> <tr> <td>accuracy</td><td>RANGE_INTERPOLATED</td></tr> </table>	lat	51.39	lon	0.54	accuracy	RANGE_INTERPOLATED																																								
lat	51.39																																														
lon	0.54																																														
accuracy	RANGE_INTERPOLATED																																														
activity	see																																														
type	landmark																																														
id	10019																																														
country	United Kingdom																																														
city	Gillingham																																														
state																																															

We can use the content field to generate an embedding, and the result will be a new document that is the same as the original, with a new field containing the vector that is the result of the embedding generation:

**Couchbase**

- [Home](#)
 - [Recipes](#)
 - [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)
-

- [Home](#)
 - [Recipes](#)
 - [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)
-



			name	Royal Engineers Museum						
			alt							
			address	Prince Arthur Road, ME4 4UG						
			directions							
			phone	+44 1634 822839						
			tollfree							
			email							
			url	http://www.remuseum.org.uk						
			hours	Tues - Fri 9.00am to 5.00pm, Sat - Sun 11.30am - 5.00pm						
			image							
			price							
			content	Adult - £6.99 for an Adult ticket that allows you to come back to history of the British Empire. A quite extensive collection of tanks and other military equipment can be seen for free.						
			geo	<table border="1"> <tr><td>lat</td><td>51.39</td></tr> <tr><td>lon</td><td>0.54</td></tr> <tr><td>accuracy</td><td>RANGE_INTERPOLATED</td></tr> </table>	lat	51.39	lon	0.54	accuracy	RANGE_INTERPOLATED
lat	51.39									
lon	0.54									
accuracy	RANGE_INTERPOLATED									
			activity	see						
			type	landmark						
			id	10019						
			country	United Kingdom						
			city	Gillingham						
			state							
			contentVector	<table border="1"> <tr><td>0</td><td>0.02</td></tr> <tr><td>...</td><td>...</td></tr> <tr><td>1535</td><td>-0.01</td></tr> </table>	0	0.02	1535	-0.01
0	0.02									
...	...									
1535	-0.01									

The resulting document is the same as the original, but with a new field `contentVector` which contains the result of embedding the content field with the active ILM. The name of the field that the embedding will be written to will default to the name of the original field with "Vector" appended. This default behavior can be overwritten with the `vectorField` flag. The resulting document is formatted with an id and content column which allows it to be piped into a `doc upsert` command to store it in the connected couchbase cluster.

```
> doc get landmark_10019 | vector enrich-doc content | doc upsert
Embedding batch 1/1
```

#	processed	success	failed	failures	cluster
0	1	1	0		local

`vector enrich-doc` can enrich more than one document at a time. To repeat what we have done above on all the landmark documents we can get all the docs using a query and pipe the result directly into the `enrich-doc` command:



Couchbase

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

' | vector enrich-doc

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



ment id as well as the contents, the same id as the original. If there select this and specify the field with

the --id-column flag. For example to use the "name" field as the id of the resulting documents do:

```
> query "SELECT name, * FROM `travel-sample` WHERE type = 'landmark'" | vector enrich-doc content --id-column name
```

4.8.4. vector enrich-text

Generates an embedding from the input text using the active_llm and outputs a json document containing the source text and the embedding.

It can be used directly on strings, which can be passed in as a positional parameter:

```
> vector enrich-text "some string" --dimension 5
Embedding batch 1/1
```

#	id	content														
0	vector-c9e02c	<table border="1"> <tr> <td>text</td><td>some string</td></tr> <tr> <td>vector</td><td> <table border="1"> <tr> <td>0</td><td>0.31</td></tr> <tr> <td>1</td><td>-0.61</td></tr> <tr> <td>2</td><td>-0.21</td></tr> <tr> <td>3</td><td>-0.59</td></tr> <tr> <td>4</td><td>-0.38</td></tr> </table> </td></tr> </table>	text	some string	vector	<table border="1"> <tr> <td>0</td><td>0.31</td></tr> <tr> <td>1</td><td>-0.61</td></tr> <tr> <td>2</td><td>-0.21</td></tr> <tr> <td>3</td><td>-0.59</td></tr> <tr> <td>4</td><td>-0.38</td></tr> </table>	0	0.31	1	-0.61	2	-0.21	3	-0.59	4	-0.38
text	some string															
vector	<table border="1"> <tr> <td>0</td><td>0.31</td></tr> <tr> <td>1</td><td>-0.61</td></tr> <tr> <td>2</td><td>-0.21</td></tr> <tr> <td>3</td><td>-0.59</td></tr> <tr> <td>4</td><td>-0.38</td></tr> </table>	0	0.31	1	-0.61	2	-0.21	3	-0.59	4	-0.38					
0	0.31															
1	-0.61															
2	-0.21															
3	-0.59															
4	-0.38															

Or piped directly in:

```
> "some string" | vector enrich-text --dimension 5
Embedding batch 1/1
```

#	id	content														
0	vector-a269a7	<table border="1"> <tr> <td>text</td><td>some string</td></tr> <tr> <td>vector</td><td> <table border="1"> <tr> <td>0</td><td>0.31</td></tr> <tr> <td>1</td><td>-0.61</td></tr> <tr> <td>2</td><td>-0.21</td></tr> <tr> <td>3</td><td>-0.59</td></tr> <tr> <td>4</td><td>-0.38</td></tr> </table> </td></tr> </table>	text	some string	vector	<table border="1"> <tr> <td>0</td><td>0.31</td></tr> <tr> <td>1</td><td>-0.61</td></tr> <tr> <td>2</td><td>-0.21</td></tr> <tr> <td>3</td><td>-0.59</td></tr> <tr> <td>4</td><td>-0.38</td></tr> </table>	0	0.31	1	-0.61	2	-0.21	3	-0.59	4	-0.38
text	some string															
vector	<table border="1"> <tr> <td>0</td><td>0.31</td></tr> <tr> <td>1</td><td>-0.61</td></tr> <tr> <td>2</td><td>-0.21</td></tr> <tr> <td>3</td><td>-0.59</td></tr> <tr> <td>4</td><td>-0.38</td></tr> </table>	0	0.31	1	-0.61	2	-0.21	3	-0.59	4	-0.38					
0	0.31															
1	-0.61															
2	-0.21															
3	-0.59															
4	-0.38															

Note that the dimension used here (5) is unrealistically low, but has just been chosen for illustrative purposes.



Couchbase

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

upsert it into the cluster:

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



0	1	1	0	local
---	---	---	---	-------

Or the output can be piped directly into `vector search` to find indexed docs with a similar vector. This makes `vector enrich-text` a very easy way to experiment with similarity search using strings. See `vector search` for an example of this.

`vector enrich-text` can also be used on larger chunks of text to produce multiple `vector docs` at once. For example say you have a text file called `some-text.txt`:

```
> open some-text.txt | vector enrich-text | doc upsert
Embedding batch 1/1
```

#	processed	success	failed	failures	cluster
0	92	92	0		local

When used on larger amounts of text `vector enrich-text` will split it into chunks of length 1024 by default, the length of the chunks can be changed with the `--chunks` flag. In this example the contents of `some-text.txt` was split into 92 chunks, resulting in 92 `vector docs` that were then upserted into the connected cluster.

Finally if you have a set of files within a directory that you want to generate `vector docs` from then you can list the files in the current directory with `ls` and pipe this list into `vector enrich-text`:

```
> ls | vector enrich-text | doc upsert
Embedding batch 1/1
```

#	processed	success	failed	failures	cluster
0	278	278	0		local

Here `vector enrich-text` will read each file, chunk the contents, retrieve the embeddings then generate the `vector docs`. Check this [recipe](https://couchbase.sh/docs/recipes.html#_simple_rag) (https://couchbase.sh/docs/recipes.html#_simple_rag) to see how `vector enrich-text` enables the implementation of simple Retrieval Augmented Generation.

**Couchbase**

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

**Flags:**

-h, --help - Display the help message for this command
 --model <String> - the chat model to ask the question

Parameters:

question <string>: the question to be asked
 context <any>: table of strings used as context for the question (optional)

`ask` can be used to ask a question of the active large language model. It can be used to answer a standalone question:

```
> ask "What is my favourite color"
I'm sorry, but as an assistant, I don't have access to personal information about you, such as your
favorite color. If you'd like to share it with me, I'd be happy to remember it for future reference!
```

Or context can be provided for the question as a list of strings. Either as a positional parameter:

```
> ask "What is my favourite color" ["My favourite color is blue"]
Your favourite color is blue.
```

This context can also be piped into the command which allows us to ask questions about docs in our database:

Charlie local in travel-sample._default._default
 > subdoc get content landmark_10019 | select content | ask "How much does it cost to go here?"
 The admission price for an adult ticket is £6.99, which allows you to come back for further visits
 within a year. There are also tickets available for children and concessionary rates. Additionally,
 there is a collection of tank-mounted bridges and other outside attractions that can be seen for
 free. Special event weekends are included in the cost of the annual ticket.

Since the context can be a list we can also use `ask` to summarize the results from various documents:

**Couchbase**

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

• [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell) ask "what activities can I see, you can explore various engineering. Activities may



1. Viewing exhibits on military equipment, vehicles, weapons, and uniforms.
2. Learning about the history of the British Empire and its impact on world events.
3. Participating in guided tours to gain more in-depth knowledge.
4. Attending special event weekends with themed activities and demonstrations.
5. Exploring the outside collection of tank-mounted bridges and other military hardware.

At the new restaurant, you can enjoy the following activities:

1. Dining on a menu that features burgers and ribs.
2. Appreciating the Hollywood-style decor and ambiance.
3. Socializing with friends or family in a lively atmosphere.
4. Trying out new dishes or drinks from the restaurant's menu.
5. Relaxing and unwinding in a cozy setting after a day of exploring or sightseeing.

The answering of questions with supplied context can be used to easily implement simple RAG.

4.10. version

The `version` command lists the version of the Couchbase shell.

```
> version
+-----+
| version | 0.92.0 |
+-----+
```

5. Reference

5.1. Config File Format

The `~/.cbsh/config` file with examples:



- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



```

default-scope = "my-scope"
default-collection = "my-collection"
# The following can be part of the config or credentials
username = "Administrator"
password = "password"
# TLS defaults to on
# tls-enabled = false
# tls-cert-path = "/path/to/cert" # either accept all certs or provide a cert path
# tls-accept-all-certs = true
# tls-validate-hostnames = false

# User display name is optional and is used to display a different name to the username in the
# prompt itself.
# This can be useful if the username that you are provided is a long randomly generated string or
# similar.
# user-display-name = "Charlie"

# Timeouts broadly apply to the operations that you would expect them to.
# That is:
# * data: commands using the kv service such as `doc`
# * query: `query` commands
# * analytics: `analytics` commands
# * search: `search` commands
# * management: commands that perform management level operations, such as `users`, `bucket`,
# `health` etc...
data-timeout = "10s"
query-timeout = "75s"
analytics-timeout = "75s"
search-timeout = "1m 15s"
management-timeout = "75s"

```

5.2. Credentials File Format

The optional `~/.cbsh/credentials` file with examples:

```

# Allows us to evolve in the future without breaking old config files
version = 1

[[cluster]]
identifier = "default"
username = "Administrator"
password = "password"
# TLS defaults to on, accepting all certs
# tls-enabled = true
# tls-cert-path = "/path/to/cert" # either accept all certs or provide a cert path
# tls-accept-all-certs = true
# tls-validate-hostnames = false

```

**Couchbase**

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



- Added support for bucket creation/management on Capella clusters
- Added support for loading sample data onto Capella clusters
- Added support for sub doc reads
- Added integration with large language models from Google, AWS and OpenAI
- Added support for named and positional query parameters
- Support document ids being ints for `doc` commands
- Support no expiry for collections
- Unset scope/collection when higher level entity changed
- Supports creation and management of Capella Clusters
- Added support for credential creation against Capella clusters
- Added support for creation of allowed CIDRs against Capella clusters
- Take username and password securely when using `cb-env register`
- Added support for managing Capella projects
- Added support for listing Capella organizations
- Enable use of `$nu` inside scripts and commands passed to cbsh

6.2. 0.75.2 - 2023-04-10

- Updated macOS build to not have any dependency on OpenSSL.
- Fixed the release workflow so that Linux release tarballs contain the `cbsh` binary.

6.3. 0.75.1 - 2023-04-13

This release contains a number of breaking changes, which are explicitly called out below. As our versioning continues to track the underlying Nushell minor version this has required breaking changes in a patch version.

- Updated config file to rename `to` (will continue to work).
- **Breaking** Updated config file to rename `hostnames` to `connstr` and changed the format to be a string.
- Added support, and detection, for different "cluster types"; Capella and Other. This allows us to modify behavior based on cluster type.
- **Breaking** Renamed `clusters health` to `health`.
- **Breaking** Renamed other `clusters ...` commands to `database ...`
- Replaced references to `cluster` with `database`.
- **Breaking** Removed support for `whoami`
- Added support for username aliases - added `display_name` to config.
- Trust the system store and Capella root CA when no certificate set.

**Couchbase**

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)

- [Home](#)
- [Recipes](#)
- [Github \(https://github.com/couchbaselabs/couchbase-shell\)](https://github.com/couchbaselabs/couchbase-shell)



- Fast fail buckets commands when used with Capella.
- Updated where config files are automatically written to.
- Fixed issue with config.nu file on Windows.
 - Pulled all beta and alpha release versions and updated version numbering
 - Couchbase Shell versions will now map to the Nushell version being used
 - Bundle Capella root CA to allow seamlessly connecting over TLS
 - Automatically detect when query_context should be sent
 - Update when SRV lookups are performed
 - Statically link OpenSSL
 - Various logging and error enhancements
 - Remove support for Capella InVpc
 - Renamed clusters managed to cb-env managed
 - Renamed clusters register/unregister to cb-env register/unregister
 - Expose CIDR in result of clusters
 - Fetch collection id over memcached rather than http

6.4. 0.75.0 - 2023-02-09

- Nushell pinned to 0.75
- Pulled all beta and alpha release versions and updated version numbering
- Couchbase Shell versions will now map to the Nushell version being used
- Bundle Capella root CA to allow seamlessly connecting over TLS
- Automatically detect when query_context should be sent
- Update when SRV lookups are performed
- Statically link OpenSSL
- Various logging and error enhancements
- Remove support for Capella InVpc
- Renamed clusters managed to cb-env managed
- Renamed clusters register/unregister to cb-env register/unregister
- Expose CIDR in result of clusters
- Fetch collection id over memcached rather than http

Last updated 2024-10-02 13:38:55 UTC